

## What is LLM as a judge?

**TL;DR:** LLM-as-a-Judge uses LLMs to evaluate AI-generated texts based on custom criteria defined in an evaluation prompt.

As you build your LLM-powered product — whether it's a chatbot, code generator, or email assistant — you need to [evaluate its quality](#).

- During development, to compare models or prompts and ensure you're improving.
- Once it's live, to monitor user interactions for quality and safety.
- Anytime you make changes, to run regression tests and ensure nothing breaks.

LLM-as-a-judge is an evaluation approach that supports all these workflows. The idea is simple: ask an LLM to "judge" the text outputs using guidelines you define.

Say, you have a chatbot. You can ask an external LLM to evaluate its responses, similar to how a human evaluator would, looking at things like:

- **Politeness:** Is the response respectful and considerate?
- **Bias:** Does the response show prejudice towards a particular group?
- **Tone:** Is the tone formal, friendly, or conversational?
- **Sentiment:** Is the emotion expressed in the text positive, negative or neutral?
- **Hallucinations:** Does this response stick to the provided context?

To apply the method, you take the text output from your AI system and feed it back into the LLM, this time alongside an **evaluation prompt**. The LLM will then return a score, label, or even a descriptive judgment — following your instructions.

The beauty of this approach is that it lets you evaluate text outputs **automatically** and look at **custom properties** specific to your use case.

For example, you can instruct the LLM to judge the helpfulness of the customer chatbot's response.